

Experiment N0. 06

Aim: To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy your first Kubernetes application.

LOs:

Theory:

This process involves installing the command-line tool for Kubernetes and using it to deploy a complete Nginx web server application.

Step 1: Install Kubectl

kubectl is the primary tool you'll use to interact with your Kubernetes cluster's API Server.

```
curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"
sudo install -o root -g root -m 0755 kubectl /usr/local/bin/kubectl
```

```
ubuntu@ip-172-31-45-31:~$ curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"
sudo install -o root -g root -m 0755 kubectl /usr/local/bin/kubectl
  % Total    % Received % Xferd  Average Speed   Time    Time     Time    Time
Time  Current          Dload  Upload   Total   Spent    Left    Speed
100 138 100 138    0     0  1469      0 --:--:-- --:--:-- --
:--:-- 1483
100 57.7M 100 57.7M    0     0  140M      0 --:--:-- --:--:-- --
:--:-- 140M
```

Step 2: Execute Kubectl Commands to Deploy an Application

We will deploy a simple Nginx application by defining its desired state in YAML manifest files.

1. **Create a Namespace:** A namespace provides a scope for names, allowing you to organize your resources.

```
kubectl create namespace my-nginx-app
```

```
ubuntu@ip-172-31-45-31:~$ kubectl create namespace my-nginx-app
namespace/my-nginx-app created
```

2. **Create a Deployment Manifest:** A **Deployment** manages your application's Pods, ensuring a specified number of replicas are always running. This command creates a **deployment.yaml** file.

```
cat <<EOF > deployment.yaml
# API version for Deployments
apiVersion: apps/v1
# Type of object
kind: Deployment
metadata:
  name: my-nginx-deployment
  # Place it in our new namespace
  namespace: my-nginx-app
spec:
  # We want 2 identical copies (Pods) of our app
  replicas: 2
  # This selector tells the Deployment which Pods to manage
  selector:
    matchLabels:
      app: nginx
  # This is the template, or blueprint, for the Pods
  template:
    metadata:
      # Pods created will have this label
      labels:
        app: nginx
    spec:
      containers:
        # Define the container to run inside the Pod
        - name: nginx-container
          image: nginx # The Docker image to use
          ports:
            - containerPort: 80 # The port the app listens on
EOF
```

```
ubuntu@ip-172-31-45-31:~$ cat <<EOF > deployment.yaml
# API version for Deployments
apiVersion: apps/v1
# Type of object
kind: Deployment
metadata:
  name: my-nginx-deployment
  # Place it in our new namespace
  namespace: my-nginx-app
spec:
  # We want 2 identical copies (Pods) of our app
  replicas: 2
  # This selector tells the Deployment which Pods to manage
```

3. **Create a Service Manifest:** A **Service** exposes your application to the network. This command creates a **service.yaml** file.

```
cat <<EOF > service.yaml
apiVersion: v1
kind: Service
metadata:
  name: my-nginx-service
  namespace: my-nginx-app
spec:
  # NodePort exposes the service on a static port on the Node's IP
  type: NodePort
  # This selector finds Pods with the label 'app: nginx'
  selector:
    app: nginx
  ports:
    # Define port mapping
    - protocol: TCP
      port: 80 # The service's internal port
      targetPort: 80 # The port on the Pod to forward traffic to
EOF
```

```
ubuntu@ip-172-31-45-31:~$ cat <<EOF > service.yaml
apiVersion: v1
kind: Service
metadata:
  name: my-nginx-service
  namespace: my-nginx-app
spec:
  # NodePort exposes the service on a static port on the Node's IP
  type: NodePort
  # This selector finds Pods with the label 'app: nginx'
  selector:
    app: nginx
  ports:
    # Define port mapping
    - protocol: TCP
      port: 80 # The service's internal port
      targetPort: 80 # The port on the Pod to forward traffic to
EOF
```

4. **Apply the Manifests:** These commands send the configuration files to the Kubernetes API Server, which then creates the resources.

```
kubectl apply -f deployment.yaml
kubectl apply -f service.yaml
```

```
ubuntu@ip-172-31-45-31:~$ kubectl apply -f deployment.yaml
kubectl apply -f service.yaml
deployment.apps/my-nginx-deployment created
service/my-nginx-service created
```

5. **Verify the Deployment:** Check the status of all the resources you created in your namespace.

`kubectl get all --namespace my-nginx-app`

```
ubuntu@ip-172-31-45-31:~$ kubectl get all --namespace my-nginx-app
NAME                                     READY   STATUS    RESTARTS   AGE
pod/my-nginx-deployment-599786d4bc-7zbcn  1/1     Running   0           15s
pod/my-nginx-deployment-599786d4bc-m6fkg  1/1     Running   0           15s

NAME                                TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)        AGE
service/my-nginx-service            NodePort    10.99.151.217 <none>        80:32018/TCP   15s

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/my-nginx-deployment  2/2     2             2           15s
```

6. **Access the Application from Your Browser:** This command forwards traffic from your EC2 instance's port 8080 to the Nginx service. This terminal must be left open.

`kubectl port-forward --namespace my-nginx-app service/my-nginx-service 8080:80 --address 0.0.0.0`

```
ubuntu@ip-172-31-45-31:~$ kubectl port-forward --namespace my-nginx-app service/my-nginx-service 8080:80 --address 0.0.0.0
Forwarding from 0.0.0.0:8080 -> 80
Handling connection for 8080
Handling connection for 8080
Handling connection for 8080
Handling connection for 8080
```

7. **Final Step: View in Your Browser**

Now, on your **local computer**, open a web browser and go to this address:

`http://<Your-EC2-Public-IP>:8080`

